

notebook_4_overdose_death_and_acs_by_census_block_merge_and_eda

October 19, 2019

1 Opioid overdose deaths and ACS dataset merge and EDA

1.0.1 Goals:

- Try normalizing death count data to town population (turn raw counts into deaths per 10k residents values) - evaluate if this improves the skew in the overdose death data
- Pull out some ACS demographics data - correlation with death counts/rate?
- EDA on ACS demographics data, some feature engineering on poverty, population, and other demographics data

1.0.2 Output:

Opioid overdose death count (by town for years 2014-2018) merged with ACS demographics data

* data/tidy_data/overdose_death_count_acs_merge.csv

pdf output in case notebook doesn't run: * products/notebook_4_overdose_death_and_acs_by_census_block_eda.pdf

```
[1]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
sns.set_style('darkgrid')
sns.set(font_scale=1.5)
```

```
[2]: death_data = pd.read_csv("../data/tidy_data/
    ↳ma_town_opioid_overdose_death_by_place_of_death_2014_to_2018.csv")
# from notebook 3 (town - census block matching)
town_block_match = pd.read_csv("../data/tidy_data/census_block_town_match.
    ↳csv")
town_w_2010pop = pd.read_csv("../data/tidy_data/
    ↳census_block_town_match_2010pop_error.csv")
# 2017 ACS dataset for MA state only, provided by Biobot
acs_17 = pd.read_csv("../data/raw_data/american_community_survey/
    ↳R12288202_SL150.csv")
```

```
[3]: death_data.head()
```

```
[3]: city_death 2014 2015 2016 2017 2018
0    abington    0    6    1    3    5
1      acton    1    2    3    0    1
2    acushnet    0    4    2    4    0
3      adams    2    3    1    0    4
4    agawam    1    2    0    4    8
```

```
[4]: # town - census block match df:
display(town_block_match.head())
display(town_w_2010pop.head())
```

```
      TOWN      GEOID10
0  wellesley  2.502140e+11
1  wellesley  2.502140e+11
2  wellesley  2.502140e+11
3  wellesley  2.502140e+11
4  wellesley  2.502140e+11
```

```
      TOWN  town_actual_2010_pop  block_est_2010_pop  count_error \
0  wellesley                27982             27982.0          0.0
1   needham                28886             28886.0          0.0
2  petersham                1234              1234.0          0.0
3   reading                24747             24747.0          0.0
4   quincy                 92271             89703.0        -2568.0
```

```
      percent_error
0          0.000000
1          0.000000
2          0.000000
3          0.000000
4          2.783106
```

```
[5]: # town mismatch between death count data and the town-census block
print(set(death_data['city_death']) - set(town_block_match['TOWN']))
print(set(town_block_match['TOWN']) - set(death_data['city_death']))
death_data['city_death'] = death_data['city_death'].str.replace('north_
→attleboro', 'north attleborough')
```

```
{'north attleboro'}
{'north attleborough'}
```

```
[6]: # ACS size:
print(acs_17.shape)
# column name format:
acs_17.columns
```

```
(4985, 2200)
```

```
[6]: Index(['Geo_FIPS', 'Geo_GEOID', 'Geo_NAME', 'Geo_QName', 'Geo_STUSAB',
        'Geo_SUMLEV', 'Geo_GEOCOMP', 'Geo_FILEID', 'Geo_LOGRECNO', 'Geo_US',
        ...,
        'SE_A10065_001', 'SE_A10065_002', 'SE_A10066_001', 'SE_A10066_002',
        'SE_A10066_003', 'SE_A10066_004', 'SE_A10066_005', 'SE_A10066_006',
        'SE_A10066_007', 'SE_A10066_008'],
        dtype='object', length=2200)
```

```
[7]: # which columns to match on?
acs_17[['Geo_FIPS', 'Geo_GEOID']].head()
# Geo_FIPS is match for GEOID10 from town - block match df
```

```
[7]:      Geo_FIPS      Geo_GEOID
0  250010101001  15000US250010101001
1  250010101002  15000US250010101002
2  250010101003  15000US250010101003
3  250010101004  15000US250010101004
4  250010101005  15000US250010101005
```

```
[8]: # mismatches between sets?
print(len(set(acs_17['Geo_FIPS']) - set(town_block_match['GEOID10'])))
print(len(set(town_block_match['GEOID10']) - set(acs_17['Geo_FIPS'])))
```

31

4

1.0.3 Potentially interesting columns to pull from ACS:

- A00002_001: Total Population
- A00002_002: Population Density (Per Sq. Mile)
- A12003_001: Civilian Population 16 to 19 Years:
- A12003_002: Not High School Graduate, Not Enrolled (Dropped Out)
- A12003_003: High School Graduate, or Enrolled (In School)
- A12002_001: Population 25 Years and Over:
- A12002_002: Less than High School
- A14006_001: Median Household Income (In 2017 Inflation Adjusted Dollars)
- A14008_001: Average Household Income
- NA- all missing - A14028_001: Gini Index
- NA - all missing - A17004_001: Total Employed Civilian Population 16 Years and Over
- NA - all missing - A17004_002: Employed Civilian Population 16 Years and Over: Agriculture, Forestry, Fishing and Hunting, and Mining
- NA - all missing - A17004_003: Employed Civilian Population 16 Years and Over: Construction
- A01001_011: 65 to 74 Years
- A01001_012: 75 to 84 Years
- A01001_013: 85 Years and Over
- NA- all missing - A13003A_001: Population Under 18 Years of Age for Whom Poverty Status Is Determined:
 - NA- all missing - A13003A_002: Living in Poverty

- NA- all missing - A13003A_003: At or Above Poverty Level
- NA- all missing -A13003B_001: Population Age 18 to 64 for Whom Poverty Status Is Determined:
 - NA- all missing - A13003B_002: Living in Poverty
 - NA- all missing - A13003B_003: At or Above Poverty Level
- NA- all missing -A13003C_001: Population Age 65 and Over for Whom Poverty Status Is Determined:
 - NA- all missing -A13003C_002: Living in Poverty
 - NA- all missing -A13003C_003: At or Above Poverty Level
- B13004_001: Population for Whom Poverty Status Is Determined:
 - B13004_002: Population for Whom Poverty Status Is Determined: Under 1.00 (Doing Poorly)
 - B13004_003: Population for Whom Poverty Status Is Determined: 1.00 to 1.99 (Struggling)
 - B13004_004: Population for Whom Poverty Status Is Determined: Under 2.00 (Poor or Struggling)
 - B13004_005: Population for Whom Poverty Status Is Determined: 2.00 and Over (Doing Ok)
- A13004_001: Population for Whom Poverty Status Is Determined:
 - A13004_002: Population for Whom Poverty Status Is Determined: Under .50
 - NA- all missing - A13004_003: Population for Whom Poverty Status Is Determined: .50 to .74
 - NA- all missing - A13004_004: Population for Whom Poverty Status Is Determined: .75 to .99
 - A13004_005: Population for Whom Poverty Status Is Determined: 1.00 to 1.49
 - A13004_006: Population for Whom Poverty Status Is Determined: 1.50 to 1.99
 - A13004_007: Population for Whom Poverty Status Is Determined: 2.00 and Over

```
[9]: acs_17_sub = acs_17[['Geo_FIPS', 'SE_A00002_001', 'SE_A00002_002',
                        'SE_A12003_001', 'SE_A12003_002',
                        'SE_A12002_001', 'SE_A12002_002', 'SE_A14006_001',
                        'SE_A14008_001',
                        # age 65+ cols:
                        'SE_A01001_011', 'SE_A01001_012', 'SE_A01001_012',
                        # poverty summarized:
                        'SE_B13004_001', 'SE_B13004_002', 'SE_B13004_003', 'SE_B13004_004',
                        ↪ 'SE_B13004_005'
                        # poverty raw?:
                        # 'SE_A13004_001', 'SE_A13004_002', 'SE_A13004_003', 'SE_A13004_004',
                        ↪ 'SE_A13004_005', 'SE_A13004_006', 'SE_A13004_007'
                        ]].copy()
```

```
[10]: acs_17_sub.head()
```

```
[10]:      Geo_FIPS  SE_A00002_001  SE_A00002_002  SE_A12003_001  SE_A12003_002  \
0  250010101001          998        116.1545           8           0
1  250010101002          314        613.6218           0           0
2  250010101003          750       3997.7830          12           0
3  250010101004          500       2019.0900           2           0
4  250010101005          390       2952.7180           0           0

      SE_A12002_001  SE_A12002_002  SE_A14006_001  SE_A14008_001  SE_A01001_011  \
0           894           46        52340.0    75538.664323          172
1           292            8        37841.0    65213.419913          107
2           638           28        58098.0    84414.854111          133
3           437           23        30396.0    46373.442623           46
4           377           28        47895.0    66060.344828           62

      SE_A01001_012  SE_A01001_012  SE_B13004_001  SE_B13004_002  SE_B13004_003  \
0           51           51          998          122          271
1           43           43          314           55           36
2           82           82          741           59           81
3           78           78          500           46          206
4           15           15          390           32           59

      SE_B13004_004  SE_B13004_005
0           393          605
1            91          223
2           140          601
3           252          248
4            91          299
```

```
[11]: # readable names:
acs_17_sub.columns = [
    'GEOID10', 'tot_pop_17', 'pop_density',
    'civ_pop_16_19', 'civ_pop_16_19_drop',
    'pop_over_25', 'pop_over_25_less_school',
    'med_house_inc', 'mean_house_inc',
    'age_65_to_74', 'age_75_to_84', 'age_85_over',
    'pop_det_poverty', 'pop_doing_poorly', 'pop_struggling',
    → 'pop_poor_or_strug', 'pop_doing_ok'
]
```

```
[12]: acs_17_sub.head()
```

```
[12]:      GEOID10  tot_pop_17  pop_density  civ_pop_16_19  civ_pop_16_19_drop  \
0  250010101001          998        116.1545           8           0
1  250010101002          314        613.6218           0           0
2  250010101003          750       3997.7830          12           0
3  250010101004          500       2019.0900           2           0
4  250010101005          390       2952.7180           0           0
```

	pop_over_25	pop_over_25_less_school	med_house_inc	mean_house_inc	\
0	894	46	52340.0	75538.664323	
1	292	8	37841.0	65213.419913	
2	638	28	58098.0	84414.854111	
3	437	23	30396.0	46373.442623	
4	377	28	47895.0	66060.344828	

	age_65_to_74	age_75_to_84	age_85_over	pop_det_poverty	pop_doing_poorly	\
0	172	51	51	998	122	
1	107	43	43	314	55	
2	133	82	82	741	59	
3	46	78	78	500	46	
4	62	15	15	390	32	

	pop_struggling	pop_poor_or_strug	pop_doing_ok
0	271	393	605
1	36	91	223
2	81	140	601
3	206	252	248
4	59	91	299

```
[13]: # calculate num and residents
acs_17_sub['over_65_count'] = acs_17_sub['age_65_to_74'] +
    acs_17_sub['age_75_to_84'] + acs_17_sub['age_85_over']
acs_17_sub.drop(['age_65_to_74', 'age_75_to_84', 'age_85_over'], axis = 1,
    inplace=True)
acs_17_sub.head()
```

```
[13]:
```

	GEOID10	tot_pop_17	pop_density	civ_pop_16_19	civ_pop_16_19_drop	\
0	250010101001	998	116.1545	8	0	
1	250010101002	314	613.6218	0	0	
2	250010101003	750	3997.7830	12	0	
3	250010101004	500	2019.0900	2	0	
4	250010101005	390	2952.7180	0	0	

	pop_over_25	pop_over_25_less_school	med_house_inc	mean_house_inc	\
0	894	46	52340.0	75538.664323	
1	292	8	37841.0	65213.419913	
2	638	28	58098.0	84414.854111	
3	437	23	30396.0	46373.442623	
4	377	28	47895.0	66060.344828	

	pop_det_poverty	pop_doing_poorly	pop_struggling	pop_poor_or_strug	\
0	998	122	271	393	
1	314	55	36	91	
2	741	59	81	140	
3	500	46	206	252	
4	390	32	59	91	

	pop_doing_ok	over_65_count
0	605	274
1	223	193
2	601	297
3	248	202
4	299	92

```
[14]: acs_17_sub.describe()
```

```
[14]:
```

	GEOID10	tot_pop_17	pop_density	civ_pop_16_19	\
count	4.985000e+03	4985.000000	4978.000000	4985.000000	
mean	2.501713e+11	1361.949649	8442.368002	76.123170	
std	7.723758e+07	670.479216	12666.048861	137.109958	
min	2.500101e+11	0.000000	0.000000	0.000000	
25%	2.500927e+11	880.000000	1082.917000	23.000000	
50%	2.501735e+11	1220.000000	3632.338000	51.000000	
75%	2.502354e+11	1696.000000	10744.957500	91.000000	
max	2.502776e+11	6760.000000	183026.000000	3499.000000	


	civ_pop_16_19_drop	pop_over_25	pop_over_25_less_school	\
count	4985.000000	4985.000000	4985.000000	
mean	2.065998	944.139619	91.891675	
std	7.653031	455.980232	101.193132	
min	0.000000	0.000000	0.000000	
25%	0.000000	619.000000	23.000000	
50%	0.000000	851.000000	58.000000	
75%	0.000000	1184.000000	128.000000	
max	110.000000	3897.000000	783.000000	

	med_house_inc	mean_house_inc	pop_det_poverty	pop_doing_poorly	\
count	4754.000000	4945.000000	4985.000000	4985.000000	
mean	82522.340766	101843.077887	1314.412638	145.947041	
std	40805.023166	52403.666245	642.014615	175.211379	
min	2499.000000	14219.285714	0.000000	0.000000	
25%	53333.000000	66678.928571	855.000000	34.000000	
50%	77321.000000	92363.461538	1184.000000	85.000000	
75%	104048.250000	123559.710145	1650.000000	191.000000	
max	250001.000000	526877.386935	4882.000000	2259.000000	

	pop_struggling	pop_poor_or_strug	pop_doing_ok	over_65_count
count	4985.000000	4985.000000	4985.000000	4985.000000
mean	165.163290	311.110331	1003.302307	240.608626
std	157.779113	289.346330	588.097455	171.729713
min	0.000000	0.000000	0.000000	0.000000
25%	53.000000	109.000000	601.000000	121.000000
50%	121.000000	225.000000	877.000000	206.000000
75%	227.000000	423.000000	1285.000000	319.000000

max 1392.000000 3260.000000 4279.000000 2178.000000

```
[15]: # code dapted from https://www.programcreek.com/python/example/96220/seaborn.
      ↪diverging_palette
acs_sub_corr = acs_17_sub.drop('GEOID10', axis=1).corr(method='spearman')
mask = np.zeros_like(acs_sub_corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
f, ax = plt.subplots(figsize=(11, 9))
cmap = sns.diverging_palette(220, 10, as_cmap=True)
sns.heatmap(acs_sub_corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
ax.set_title('Correlation matrix\nACS 2017 data')
plt.show()
```



Notes: * A lot of features are pretty correlated * poverty counts and mean/med income have strong correlations * Will wait to do any kind of summarizing after the merge with town/census block df


```
[16]: print(acs_17_sub.columns)
      print(town_block_match.columns)
```

```
Index(['GEOID10', 'tot_pop_17', 'pop_density', 'civ_pop_16_19',
      'civ_pop_16_19_drop', 'pop_over_25', 'pop_over_25_less_school',
      'med_house_inc', 'mean_house_inc', 'pop_det_poverty',
      'pop_doing_poorly', 'pop_struggling', 'pop_poor_or_strug',
      'pop_doing_ok', 'over_65_count'],
      dtype='object')
Index(['TOWN', 'GEOID10'], dtype='object')
```

```
[17]: town_block = town_block_match.merge(town_w_2010pop, on='TOWN', how='inner')
      print(town_block.shape)
      print(len(set(town_block['TOWN'])))
      print(len(set(town_block['GEOID10'])))
      town_block.columns = ['city_death'] + list(town_block.columns)[1:]
      # result:
      town_block.head()
```

```
(4958, 6)
351
4958
```

```
[17]: city_death      GEOID10  town_actual_2010_pop  block_est_2010_pop  \
0  wellesley  2.502140e+11          27982          27982.0
1  wellesley  2.502140e+11          27982          27982.0
2  wellesley  2.502140e+11          27982          27982.0
3  wellesley  2.502140e+11          27982          27982.0
4  wellesley  2.502140e+11          27982          27982.0

      count_error  percent_error
0              0.0              0.0
1              0.0              0.0
2              0.0              0.0
3              0.0              0.0
4              0.0              0.0
```

```
[18]: town_acs_merge = town_block.merge(acs_17_sub, on='GEOID10', how='inner')
      print(town_acs_merge.shape)
      print(town_acs_merge.columns)
      town_acs_merge.head()
```

```
(4954, 20)
Index(['city_death', 'GEOID10', 'town_actual_2010_pop', 'block_est_2010_pop',
      'count_error', 'percent_error', 'tot_pop_17', 'pop_density',
      'civ_pop_16_19', 'civ_pop_16_19_drop', 'pop_over_25',
      'pop_over_25_less_school', 'med_house_inc', 'mean_house_inc',
```

```

    'pop_det_poverty', 'pop_doing_poorly', 'pop_struggling',
    'pop_poor_or_strug', 'pop_doing_ok', 'over_65_count'],
    dtype='object')

```

```

[18]: city_death      GEOID10  town_actual_2010_pop  block_est_2010_pop  \
0  wellesley  2.502140e+11          27982          27982.0
1  wellesley  2.502140e+11          27982          27982.0
2  wellesley  2.502140e+11          27982          27982.0
3  wellesley  2.502140e+11          27982          27982.0
4  wellesley  2.502140e+11          27982          27982.0

    count_error  percent_error  tot_pop_17  pop_density  civ_pop_16_19  \
0           0.0           0.0         1101    4438.051           96
1           0.0           0.0           924    2064.697          108
2           0.0           0.0           881    1947.536           9
3           0.0           0.0          1177    2961.899          38
4           0.0           0.0           767    3843.039          10

    civ_pop_16_19_drop  pop_over_25  pop_over_25_less_school  med_house_inc  \
0                   0           636                   0          183879.0
1                   0           634                   0          250001.0
2                   0           552                  11          181786.0
3                   0           685                  30          129071.0
4                   0           590                  22           86827.0

    mean_house_inc  pop_det_poverty  pop_doing_poorly  pop_struggling  \
0  214803.395062          1101           0           0
1  505804.304636           924           0           0
2  244652.447552           875          12           9
3  240321.641791          1177          144          100
4  145115.151515           767           64           95

    pop_poor_or_strug  pop_doing_ok  over_65_count
0                   0          1101          137
1                   0           924          197
2                  21           854          231
3                 244           933          105
4                 159           608          275

```

To summarize, some columns need to be added by town (count columns) and some need to be averaged

First, columns that will be added up by group:

```

[19]: sum_stat_cols = [
    'city_death', 'tot_pop_17', 'over_65_count',
    'civ_pop_16_19', 'civ_pop_16_19_drop',
    'pop_over_25', 'pop_over_25_less_school',
    'pop_det_poverty', 'pop_doing_poorly', 'pop_struggling',
    'pop_poor_or_strug', 'pop_doing_ok'
]

```

```
]
town_17_pop = town_acs_merge.groupby('city_death').sum().
    ↪reset_index()[sum_stat_cols]
town_17_pop.head()
```

```
[19]:  city_death  tot_pop_17  over_65_count  civ_pop_16_19  civ_pop_16_19_drop  \
0    abington    16275        2469          753          26
1      acton    23455        4001         1476          18
2  acushnet    10443        2431          564          62
3      adams     8211        1764          333          23
4    agawam    27769        6195         1177           0

    pop_over_25  pop_over_25_less_school  pop_det_poverty  pop_doing_poorly  \
0         11377                615          16194          579
1         16161                397          23307          893
2          7635               1397          10336          422
3          6095                723           8209          910
4         20674               1602          26925         2553

    pop_struggling  pop_poor_or_strug  pop_doing_ok
0             1626                2205          13989
1              973                1866          21441
2             1844                2266           8070
3             1187                2097           6112
4             3841                6394          20531
```

Columns that will be averaged by group: Note: for 2010 actual and estimated error columns, as well as the count and percent error - this will return the values from the original imported df

```
[20]: mean_stat_cols = ['city_death', 'town_actual_2010_pop', 'block_est_2010_pop',
    'count_error', 'percent_error', 'pop_density',
    'med_house_inc', 'mean_house_inc']
town_stats = town_acs_merge.groupby('city_death').mean().
    ↪reset_index()[mean_stat_cols]
town_stats.head()
```

```
[20]:  city_death  town_actual_2010_pop  block_est_2010_pop  count_error  \
0    abington          15985.0          15985.0           0.0
1      acton          21924.0          21924.0           0.0
2  acushnet          10303.0          10303.0           0.0
3      adams           8485.0           8485.0           0.0
4    agawam          28438.0          27621.0         -817.0

    percent_error  pop_density  med_house_inc  mean_house_inc
0      0.000000  1932.969130   87156.000000   98809.035505
1      0.000000  1257.583593  139890.466667  156680.203867
2      0.000000  1152.357871   69624.714286   80333.175842
3      0.000000  1982.318840   48445.400000   60968.594660
4      2.872917  1897.273569   65490.125000   79464.234446
```

```
[21]: town_merge = town_17_pop.merge(town_stats, on='city_death', how='inner')
print(town_merge.shape)
town_merge.head()
```

(347, 19)

```
[21]:  city_death  tot_pop_17  over_65_count  civ_pop_16_19  civ_pop_16_19_drop  \
0  abington      16275      2469          753          26
1    acton      23455      4001         1476          18
2  acushnet     10443      2431          564          62
3    adams       8211      1764          333          23
4   agawam     27769      6195         1177           0

    pop_over_25  pop_over_25_less_school  pop_det_poverty  pop_doing_poorly  \
0        11377              615          16194          579
1        16161              397          23307          893
2         7635             1397          10336          422
3         6095              723           8209          910
4        20674             1602          26925         2553

    pop_struggling  pop_poor_or_strug  pop_doing_ok  town_actual_2010_pop  \
0          1626             2205          13989          15985.0
1           973             1866          21441          21924.0
2          1844             2266           8070          10303.0
3          1187             2097           6112           8485.0
4          3841             6394          20531          28438.0

    block_est_2010_pop  count_error  percent_error  pop_density  med_house_inc  \
0          15985.0           0.0      0.000000  1932.969130  87156.000000
1          21924.0           0.0      0.000000  1257.583593  139890.466667
2          10303.0           0.0      0.000000  1152.357871  69624.714286
3           8485.0           0.0      0.000000  1982.318840  48445.400000
4          27621.0          -817.0      2.872917  1897.273569  65490.125000

    mean_house_inc
0    98809.035505
1   156680.203867
2    80333.175842
3    60968.594660
4    79464.234446
```

Some summary stats: * drop out rate among 16-19 year olds * estimate of proportion of population that's over 25 with less than a high school education

```
[22]: town_merge['drop_out'] = (town_merge['civ_pop_16_19_drop'] * 100) /_
      ->town_merge['civ_pop_16_19']
town_merge['less_than_hs_ed'] = (town_merge['pop_over_25_less_school'] * 100) /_
      ->town_merge['pop_over_25']
```

```
town_merge.drop(['pop_over_25', 'pop_over_25_less_school',
→'civ_pop_16_19_drop', 'civ_pop_16_19'], axis = 1, inplace=True)
town_merge.head()
```

```
[22]:  city_death  tot_pop_17  over_65_count  pop_det_poverty  pop_doing_poorly \
0    abington    16275      2469          16194          579
1      acton    23455      4001          23307          893
2    acushnet    10443      2431          10336          422
3      adams     8211      1764           8209          910
4    agawam    27769      6195          26925         2553

      pop_struggling  pop_poor_or_strug  pop_doing_ok  town_actual_2010_pop \
0           1626          2205          13989         15985.0
1           973          1866          21441         21924.0
2          1844          2266           8070         10303.0
3          1187          2097           6112          8485.0
4          3841          6394          20531         28438.0

      block_est_2010_pop  count_error  percent_error  pop_density  med_house_inc \
0          15985.0          0.0          0.000000  1932.969130  87156.000000
1          21924.0          0.0          0.000000  1257.583593  139890.466667
2          10303.0          0.0          0.000000  1152.357871  69624.714286
3           8485.0          0.0          0.000000  1982.318840  48445.400000
4         27621.0         -817.0          2.872917  1897.273569  65490.125000


      mean_house_inc  drop_out  less_than_hs_ed
0    98809.035505    3.452855      5.405643
1   156680.203867    1.219512      2.456531
2    80333.175842   10.992908     18.297315
3    60968.594660    6.906907     11.862182
4   79464.234446    0.000000      7.748863
```

```
[23]: # before calculated dropout by block - had many zeroes - is this alternative
→method more informative?
town_merge[['drop_out', 'less_than_hs_ed']].describe()
```

```
[23]:      drop_out  less_than_hs_ed
count  346.000000    347.000000
mean    2.355802     6.544580
std     4.012968     4.750990
min     0.000000     0.000000
25%     0.000000     3.235427
50%     0.318954     5.405643
75%     3.184586     8.318470
max     30.000000    32.336132
```

```
[24]: # what's the distribution of these 2 new variables and what is the relationship
→between them?
sns.pairplot(town_merge[['drop_out', 'less_than_hs_ed']].dropna())
```

```
plt.show()
```



New dropout variable has a lot of zeroes - probably not useful - drop it

```
[25]: town_merge.drop('drop_out', axis=1, inplace=True)
      town_merge.head()
```

```
[25]:  city_death  tot_pop_17  over_65_count  pop_det_poverty  pop_doing_poorly  \
0    abington    16275    2469    16194    579
1      acton    23455    4001    23307    893
2  acushnet    10443    2431    10336    422
3      adams     8211    1764     8209    910
4    agawam    27769    6195    26925   2553


      pop_struggling  pop_poor_or_strug  pop_doing_ok  town_actual_2010_pop  \
0           1626           2205           13989    15985.0
1            973           1866           21441    21924.0
2           1844           2266            8070    10303.0
3           1187           2097            6112     8485.0
4           3841           6394           20531    28438.0
```

	block_est_2010_pop	count_error	percent_error	pop_density	med_house_inc	\
0	15985.0	0.0	0.000000	1932.969130	87156.000000	
1	21924.0	0.0	0.000000	1257.583593	139890.466667	
2	10303.0	0.0	0.000000	1152.357871	69624.714286	
3	8485.0	0.0	0.000000	1982.318840	48445.400000	
4	27621.0	-817.0	2.872917	1897.273569	65490.125000	

	mean_house_inc	less_than_hs_ed
0	98809.035505	5.405643
1	156680.203867	2.456531
2	80333.175842	18.297315
3	60968.594660	11.862182
4	79464.234446	7.748863

Maybe the proportion of people over 65 in a town could be useful?


```
[26]: town_merge['over_65_prop'] = town_merge['over_65_count'] / town_merge['tot_pop_17']
      town_merge['over_65_prop'].hist()
      plt.xlabel('Proportion of Population over 65')
      plt.show()
```




The different columns relating to poverty need to be summarized in some way.

What the different ratios of income/poverty mean: * Doing Poorly: Under 1.00 * Struggling: 1.00 to 1.99 * Poor or Struggling: Under 2.00 * Doing OK: 2.00 and Over


```
[27]: # poverty calc
town_merge.head()
town_merge['at_or_below_pov_prop'] = town_merge['pop_doing_poorly'] /  $\square$ 
       $\rightarrow$ town_merge['pop_det_poverty']
town_merge['at_or_below_pov_prop'].hist(bins=10)
plt.xlabel('Proportion of Residents with Income/Poverty < 1.00')
plt.show()
```



```
[28]: town_merge['pop_struggling_prop'] = town_merge['pop_struggling'] /  $\square$ 
       $\rightarrow$ town_merge['pop_det_poverty']
town_merge['pop_struggling_prop'].hist(bins=10)
plt.xlabel('Proportion of Residents Struggling (1.00 < Income/Poverty < 2.00)')
plt.show()
```

```
[29]: # what's the relationship between these 2?
plt.figure(figsize=(14,10))
sns.regplot(x='at_or_below_pov_prop', y='pop_struggling_prop', data=town_merge)
plt.xlabel('Population at or below Inc/Pov < 1.0 (Doing Poorly)')
plt.ylabel('Population 1.0 < Inc/Pov < 2.0 (Struggling)')
plt.show()
```



```
[30]: # drop count data
town_merge.drop(['pop_det_poverty', 'pop_doing_poorly', 'pop_struggling',
                → 'pop_poor_or_strug', 'pop_doing_ok'], axis=1, inplace=True)
town_merge.head()
```

```
[30]:  city_death  tot_pop_17  over_65_count  town_actual_2010_pop  \
0    abington    16275      2469          15985.0
1      acton    23455      4001          21924.0
2  acushnet    10443      2431          10303.0
3      adams     8211      1764           8485.0
4    agawam    27769      6195         28438.0

    block_est_2010_pop  count_error  percent_error  pop_density  med_house_inc  \
0          15985.0          0.0      0.000000  1932.969130  87156.000000
1          21924.0          0.0      0.000000  1257.583593  139890.466667
2          10303.0          0.0      0.000000  1152.357871   69624.714286
3           8485.0          0.0      0.000000  1982.318840   48445.400000
4          27621.0         -817.0      2.872917  1897.273569   65490.125000

    mean_house_inc  less_than_hs_ed  over_65_prop  at_or_below_pov_prop  \
0   98809.035505      5.405643      0.151705          0.035754
1  156680.203867      2.456531      0.170582          0.038315
2   80333.175842     18.297315      0.232788          0.040828
3   60968.594660     11.862182      0.214834          0.110854
4   79464.234446      7.748863      0.223090          0.094819

    pop_struggling_prop
0          0.100408
1          0.041747
2          0.178406
3          0.144597
4          0.142656
```

```
[31]: # distributions of current
town_merge.describe()
```

```
[31]:  tot_pop_17  over_65_count  town_actual_2010_pop  block_est_2010_pop  \
count      347.000000      347.000000          347.000000          347.000000
mean    19490.746398    3434.824207      18858.389049      18794.827089
std     41540.511381    5642.967441      39009.479522      38920.193487
min         34.000000      15.000000        75.000000        75.000000
25%      4196.000000      781.000000       4008.000000       4008.000000
50%     10560.000000     1977.000000      10300.000000      10209.000000
75%      22704.000000     4385.000000      21691.500000      21691.500000
max     668541.000000    85040.000000     617594.000000     616852.000000

    count_error  percent_error  pop_density  med_house_inc  \
```

count	347.000000	347.000000	347.000000	346.000000
mean	-63.561960	5.979779	2292.685930	88305.727278
std	445.473371	44.749992	4049.633320	29088.836206
min	-3861.000000	0.000000	2.578370	38909.750000
25%	0.000000	0.000000	285.202885	68215.428571
50%	0.000000	0.000000	933.183133	83124.625000
75%	0.000000	0.000000	2452.889962	102446.289216
max	1323.000000	733.532934	30236.970333	203026.750000

	mean_house_inc	less_than_hs_ed	over_65_prop	at_or_below_pov_prop \
count	347.000000	347.000000	347.000000	347.000000
mean	109444.035932	6.544580	0.207933	0.073728
std	39888.981321	4.750990	0.068754	0.047828
min	50750.537570	0.000000	0.086886	0.005354
25%	82839.316239	3.235427	0.166202	0.042404
50%	100209.367399	5.405643	0.193279	0.059730
75%	126078.167762	8.318470	0.230995	0.092880
max	316351.858774	32.336132	0.532847	0.332260

	pop_struggling_prop
count	347.000000
mean	0.104279
std	0.054009
min	0.000000
25%	0.064869
50%	0.092496
75%	0.132362
max	0.304869

```
[32]: print(set(death_data['city_death']) - set(town_merge['city_death']))
      print(set(town_merge['city_death']) - set(death_data['city_death']))
```

```
{'monroe', 'egremont', 'charlemont', 'worthington'}
set()
```

```
[33]: # combine town info + opioid overdose death count data
      full_merge = town_merge.merge(death_data, on='city_death', how='inner')
      print(full_merge.shape)
      full_merge.head()
```

```
(347, 19)
```

```
[33]: city_death  tot_pop_17  over_65_count  town_actual_2010_pop \
0  abington      16275      2469      15985.0
1  acton         23455      4001      21924.0
2  acushnet      10443      2431      10303.0
3  adams         8211      1764      8485.0
```

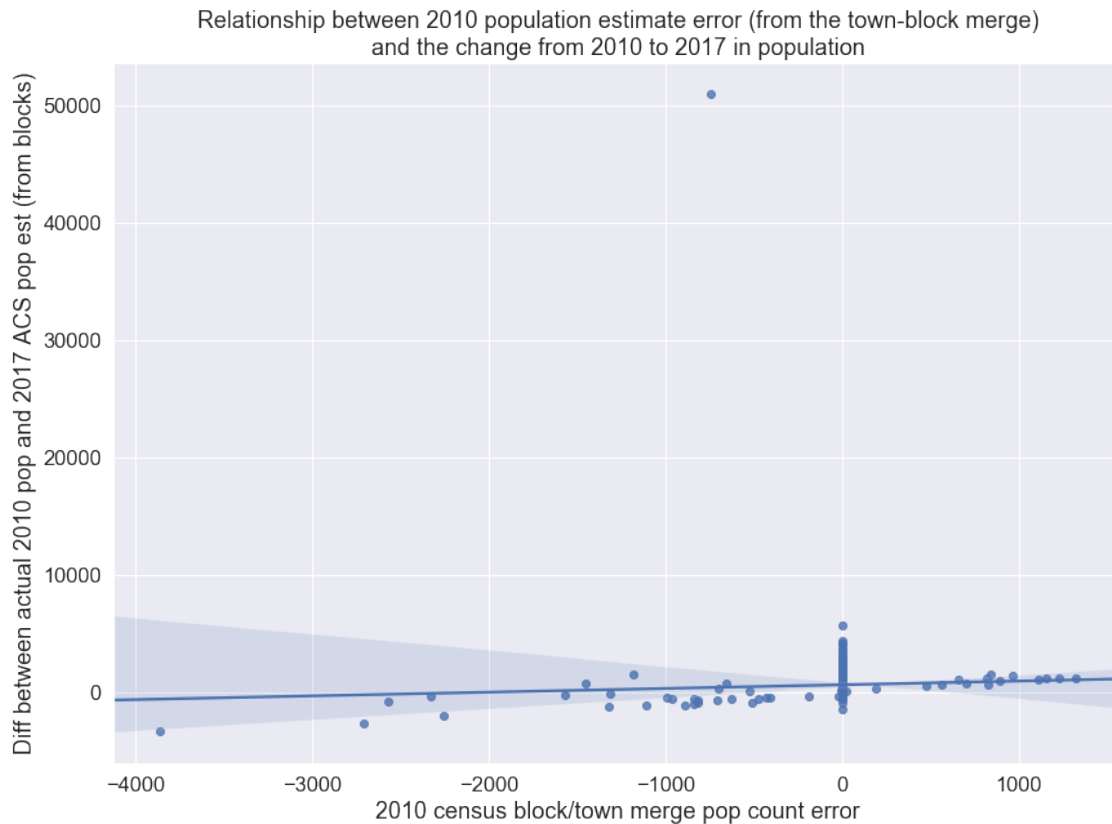
4	agawam	27769	6195	28438.0
---	--------	-------	------	---------

	block_est_2010_pop	count_error	percent_error	pop_density	med_house_inc \
0	15985.0	0.0	0.000000	1932.969130	87156.000000
1	21924.0	0.0	0.000000	1257.583593	139890.466667
2	10303.0	0.0	0.000000	1152.357871	69624.714286
3	8485.0	0.0	0.000000	1982.318840	48445.400000
4	27621.0	-817.0	2.872917	1897.273569	65490.125000

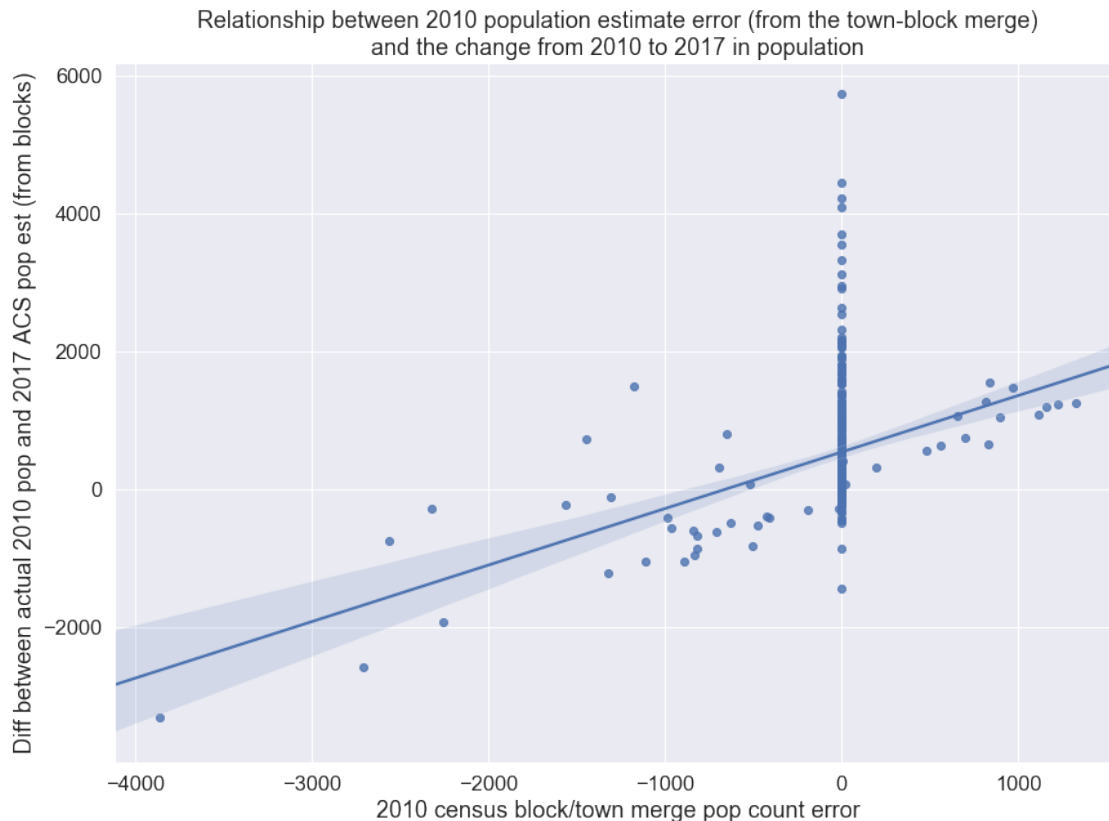
	mean_house_inc	less_than_hs_ed	over_65_prop	at_or_below_pov_prop \
0	98809.035505	5.405643	0.151705	0.035754
1	156680.203867	2.456531	0.170582	0.038315
2	80333.175842	18.297315	0.232788	0.040828
3	60968.594660	11.862182	0.214834	0.110854
4	79464.234446	7.748863	0.223090	0.094819

	pop_struggling_prop	2014	2015	2016	2017	2018
0	0.100408	0	6	1	3	5
1	0.041747	1	2	3	0	1
2	0.178406	0	4	2	4	0
3	0.144597	2	3	1	0	4
4	0.142656	1	2	0	4	8

```
[34]: full_merge['pop_change_10_to_17'] = full_merge['tot_pop_17'] -
      ↪full_merge['town_actual_2010_pop']
plt.figure(figsize=(14,10))
sns.regplot(x='count_error', y='pop_change_10_to_17', data=full_merge)
plt.xlabel('2010 census block/town merge pop count error')
plt.ylabel('Diff between actual 2010 pop and 2017 ACS pop est (from blocks)')
plt.title('Relationship between 2010 population estimate error (from the
      ↪town-block merge)\nand the change from 2010 to 2017 in population')
plt.show()
```

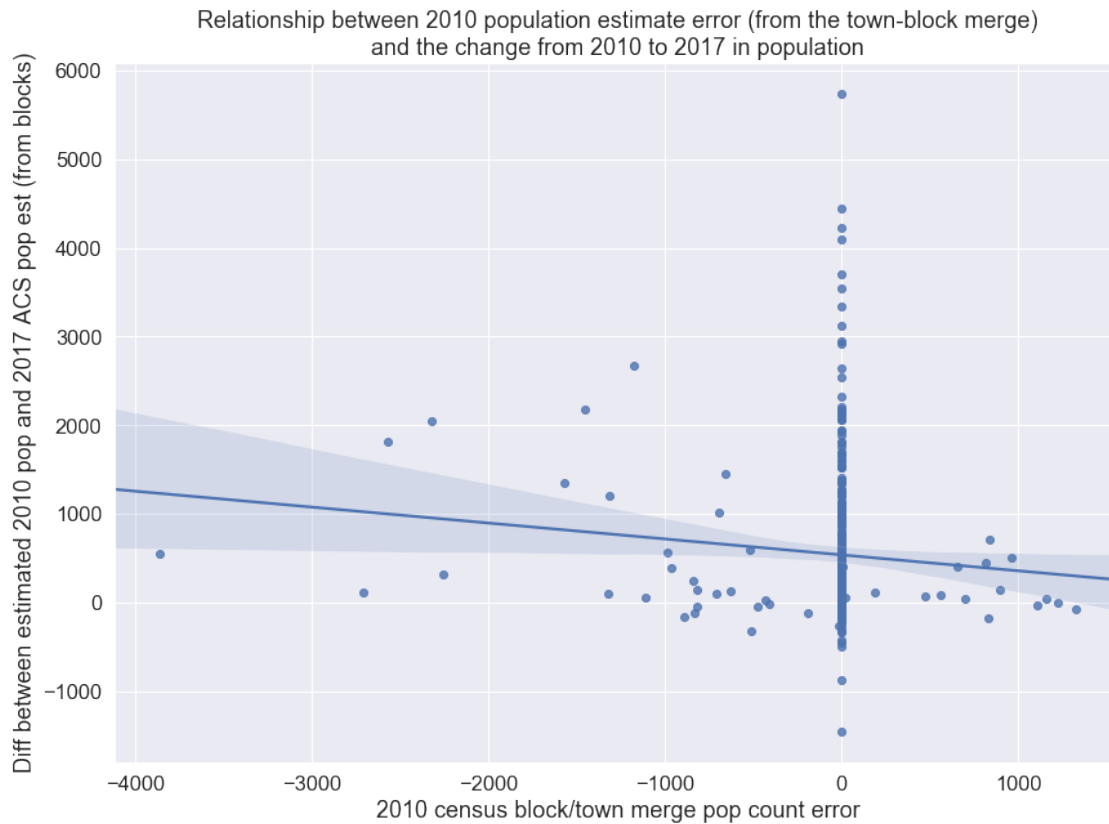


```
[35]: # take out the one big change point (Boston)
plt.figure(figsize=(14,10))
sns.regplot(x='count_error', y='pop_change_10_to_17',
            data=full_merge[full_merge['pop_change_10_to_17'] < 40000])
plt.xlabel('2010 census block/town merge pop count error')
plt.ylabel('Diff between actual 2010 pop and 2017 ACS pop est (from blocks)')
plt.title('Relationship between 2010 population estimate error (from the_
→town-block merge)\nand the change from 2010 to 2017 in population')
plt.show()
```



Notes: * There does seem to be some relationship between the error in the 2010 estimation and the difference between the 2017 population estimate (also derived from census block data) and the 2010 actual count - makes sense * But it's not perfectly 1-to-1 - the difference between (2017 estimate - 2010 actual) and (2010 estimate - 2010 actual) seems to be smaller

```
[36]: full_merge['pop_change_10_to_17_est'] = full_merge['tot_pop_17'] -
      →full_merge['block_est_2010_pop']
plt.figure(figsize=(14,10))
sns.regplot(x='count_error', y='pop_change_10_to_17_est',
            data=full_merge[full_merge['pop_change_10_to_17_est'] < 40000])
plt.xlabel('2010 census block/town merge pop count error')
plt.ylabel('Diff between estimated 2010 pop and 2017 ACS pop est (from blocks)')
plt.title('Relationship between 2010 population estimate error (from the
      →town-block merge)\nand the change from 2010 to 2017 in population')
plt.show()
```




```
[37]: print(sum(full_merge['pop_change_10_to_17_est'] == 0))
      print(sum(full_merge['pop_change_10_to_17'] == 0))
```

0
0

Will use 2010 estimate derived from the town - census block merge vs 2017 estimate derived from merge to estimate growth rate (hopefully the error made from census blocks-town assignment means a similar error for both estimates)

```
[38]: full_merge['pop_change_rate'] = full_merge['pop_change_10_to_17_est'] / (2017 -
      ↪ 2010)
      full_merge['pop_change_rate'].hist(bins=20)
      plt.xlabel('Calculated population change rate per year (2010 to 2017)')
      plt.show()
```



```
[39]: full_merge[full_merge['pop_change_rate'] > 1000]
```

```
[39]:  city_death  tot_pop_17  over_65_count  town_actual_2010_pop  \
35      boston      668541          85040          617594.0

      block_est_2010_pop  count_error  percent_error  pop_density  \
35          616852.0        -742.0         0.120144  27786.891612


      med_house_inc  mean_house_inc  ...  at_or_below_pov_prop  \
35  71196.049505    92167.940523  ...              0.205154

      pop_struggling_prop  2014  2015  2016  2017  2018  pop_change_10_to_17  \
35          0.158644    167    226    259    279    245          50947.0

      pop_change_10_to_17_est  pop_change_rate
35          51689.0          7384.142857

[1 rows x 22 columns]
```

```
[40]: # remove extreme point (Boston)
full_merge[full_merge['pop_change_rate'] < 1000]['pop_change_rate'].
    hist(bins=20)
plt.xlabel('Calculated population change rate per year (2010 to 2017)')
plt.show()
```

```
[41]: full_merge[(full_merge['pop_change_rate'] > -10) &
      →(full_merge['pop_change_rate'] < 10)]
```

```
[41]:
```

	city_death	tot_pop_17	over_65_count	town_actual_2010_pop	\
29	bernardston	1074	261	2129.0	
33	blandford	1259	306	1233.0	
36	bourne	19832	4614	19754.0	
41	brewster	9856	3274	9820.0	
45	brookfield	3406	791	3390.0	
..	
307	warren	5199	554	5135.0	
309	washington	499	158	538.0	
315	wendell	864	204	848.0	
336	williamsburg	2481	578	2482.0	
341	windsor	909	184	899.0	

	block_est_2010_pop	count_error	percent_error	pop_density	\
29	1018.0	-1111.0	52.184124	86.640560	
33	1233.0	0.0	0.000000	24.411060	
36	19776.0	22.0	0.111370	1009.739420	
41	9820.0	0.0	0.000000	491.873789	
45	3390.0	0.0	0.000000	237.359100	
..	
307	5135.0	0.0	0.000000	279.391975	
309	538.0	0.0	0.000000	13.137440	
315	848.0	0.0	0.000000	27.134650	
336	2482.0	0.0	0.000000	119.907030	

```
341          899.0          0.0          0.000000          25.976800
```

```

    med_house_inc  mean_house_inc  ...  at_or_below_pov_prop  \
29    70500.000000    82633.333333  ...              0.035382
33    62875.000000    78223.844732  ...              0.070691
36    69166.937500    91659.357901  ...              0.068506
41    71159.111111    93807.274446  ...              0.048019
45    63971.333333    86589.528959  ...              0.091930
..          ...          ...          ...
307   68332.000000    70051.858136  ...              0.112906
309   86389.000000    91221.120690  ...              0.028169
315   42750.000000    65966.497462  ...              0.175581
336   74244.500000    95587.601080  ...              0.108021
341   81875.000000   110717.195767  ...              0.055006
```

```

    pop_struggling_prop  2014  2015  2016  2017  2018  pop_change_10_to_17  \
29              0.220670      0      0      0      0      0             -1055.0
33              0.155679      0      0      0      0      0              26.0
36              0.112949      2      4      6      1      4              78.0
41              0.122623      0      1      1      1      0              36.0
45              0.110553      0      0      0      2      1              16.0
..              ...      ...      ...      ...      ...      ...
307             0.205424      1      0      0      0      1              64.0
309             0.072435      0      0      0      0      0             -39.0
315             0.172093      0      0      0      0      0              16.0
336             0.077791      0      0      0      0      1              -1.0
341             0.063806      0      0      0      0      0              10.0
```

```

    pop_change_10_to_17_est  pop_change_rate
29              56.0              8.000000
33              26.0              3.714286
36              56.0              8.000000
41              36.0              5.142857
45              16.0              2.285714
..              ...              ...
307             64.0              9.142857
309            -39.0             -5.571429
315             16.0              2.285714
336             -1.0             -0.142857
341             10.0              1.428571
```

```
[62 rows x 22 columns]
```

```
[42]: # number of towns with negative growth rate
full_merge[full_merge['pop_change_rate'] < 0].shape
```

```
[42]: (60, 22)
```

- Tried to pull in more years of ACS data, but it was challenging to find and extract the data
- Will use rough estimates to guess population changes

```
[43]: # rough estimates of yearly population
full_merge['tot_pop_16'] = full_merge['tot_pop_17'] -
    ↳full_merge['pop_change_rate']
full_merge['tot_pop_15'] = full_merge['tot_pop_17'] -
    ↳full_merge['pop_change_rate'] * 2
full_merge['tot_pop_14'] = full_merge['tot_pop_17'] -
    ↳full_merge['pop_change_rate'] * 3
full_merge['tot_pop_13'] = full_merge['tot_pop_17'] -
    ↳full_merge['pop_change_rate'] * 4
full_merge.head()
```

```
[43]: city_death  tot_pop_17  over_65_count  town_actual_2010_pop  \
0    abington      16275           2469           15985.0
1      acton      23455           4001           21924.0
2    acushnet      10443           2431           10303.0
3      adams       8211           1764            8485.0
4    agawam      27769           6195           28438.0

    block_est_2010_pop  count_error  percent_error  pop_density  med_house_inc  \
0           15985.0           0.0           0.000000  1932.969130  87156.000000
1           21924.0           0.0           0.000000  1257.583593  139890.466667
2           10303.0           0.0           0.000000  1152.357871   69624.714286
3            8485.0           0.0           0.000000  1982.318840   48445.400000
4          27621.0          -817.0           2.872917  1897.273569   65490.125000

    mean_house_inc  ...  2016  2017  2018  pop_change_10_to_17  \
0    98809.035505  ...    1    3    5           290.0
1   156680.203867  ...    3    0    1          1531.0
2    80333.175842  ...    2    4    0           140.0
3    60968.594660  ...    1    0    4          -274.0
4    79464.234446  ...    0    4    8          -669.0

    pop_change_10_to_17_est  pop_change_rate  tot_pop_16  tot_pop_15  \
0              290.0           41.428571  16233.571429  16192.142857
1             1531.0          218.714286  23236.285714  23017.571429
2              140.0           20.000000  10423.000000  10403.000000
3             -274.0          -39.142857   8250.142857   8289.285714
4              148.0           21.142857  27747.857143  27726.714286

    tot_pop_14  tot_pop_13
0  16150.714286  16109.285714
1  22798.857143  22580.142857
2  10383.000000  10363.000000
3   8328.428571   8367.571429
4  27705.571429  27684.428571
```

[5 rows x 26 columns]

```
[44]: # classify a town as shrinking or growing based on the rate of population change
full_merge['town_status'] = ['grown' if x > 10 else 'shrunk' if x < -10 else
    → 'uncertain' for x in full_merge['pop_change_rate']]
full_merge['town_status'].value_counts()
```

```
[44]: grown      252
      uncertain   62
      shrunk     33
      Name: town_status, dtype: int64
```

```
[45]: # simple definition of urban vs rural
full_merge['urb_v_rur'] = ['urban' if x >= 50000 else 'rural' for x in
    → full_merge['tot_pop_17']]
full_merge['urb_v_rur'].value_counts()
```

```
[45]: rural      322
      urban     25
      Name: urb_v_rur, dtype: int64
```

I want to experiment with converting the death count to a death rate by normalizing it to the population estimate

- will it change the distribution? (the count data is very skewed)
- what is the relationship between the death rate and some of the other variables pulled out from the 2017 ACS?

```
[46]: # calculate opioid overdose death rate per 10k residents (so that numbers
    → aren't just small decimals)
rate_exp = full_merge.copy()
rate_exp['death_rate_17'] = (rate_exp['2017'] / rate_exp['tot_pop_17']) * 10000
rate_exp.head()
```

```
[46]:  city_death  tot_pop_17  over_65_count  town_actual_2010_pop  \
0    abington    16275         2469          15985.0
1      acton    23455         4001          21924.0
2   acushnet    10443         2431          10303.0
3      adams     8211         1764           8485.0
4    agawam    27769         6195          28438.0

      block_est_2010_pop  count_error  percent_error  pop_density  med_house_inc  \
0           15985.0         0.0         0.000000    1932.969130    87156.000000
1           21924.0         0.0         0.000000    1257.583593    139890.466667
2           10303.0         0.0         0.000000    1152.357871    69624.714286
3            8485.0         0.0         0.000000    1982.318840    48445.400000
4          27621.0        -817.0         2.872917    1897.273569    65490.125000

      mean_house_inc  ...  pop_change_10_to_17  pop_change_10_to_17_est  \
0    98809.035505  ...              290.0              290.0
```

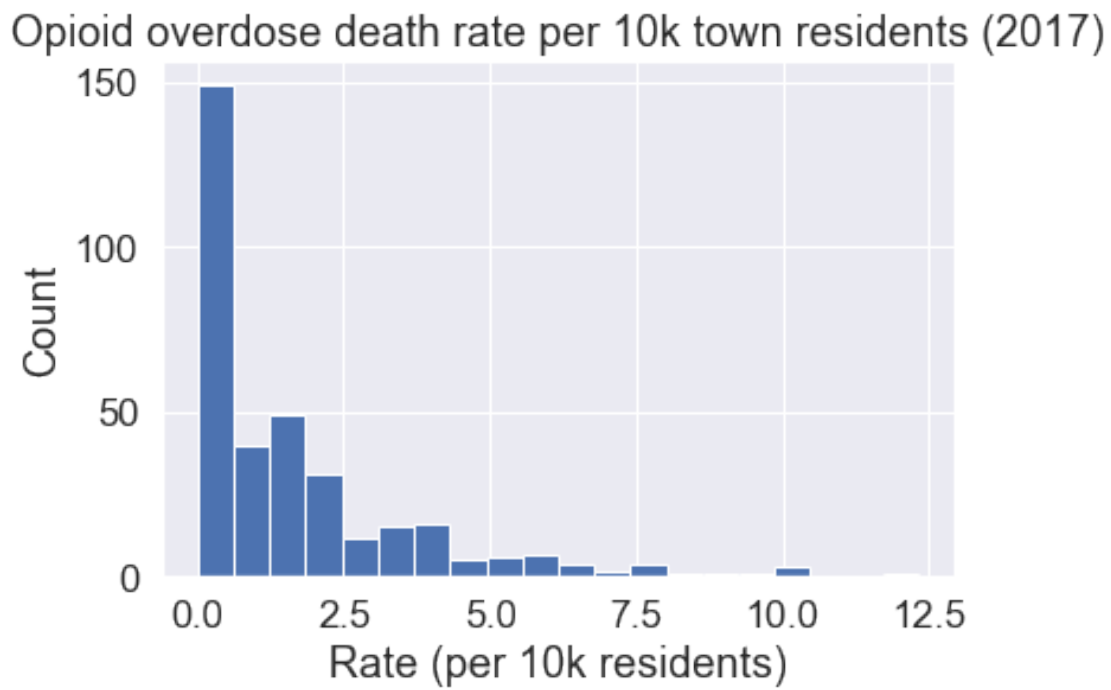
1	156680.203867	...	1531.0	1531.0
2	80333.175842	...	140.0	140.0
3	60968.594660	...	-274.0	-274.0
4	79464.234446	...	-669.0	148.0

	pop_change_rate	tot_pop_16	tot_pop_15	tot_pop_14	tot_pop_13 \
0	41.428571	16233.571429	16192.142857	16150.714286	16109.285714
1	218.714286	23236.285714	23017.571429	22798.857143	22580.142857
2	20.000000	10423.000000	10403.000000	10383.000000	10363.000000
3	-39.142857	8250.142857	8289.285714	8328.428571	8367.571429
4	21.142857	27747.857143	27726.714286	27705.571429	27684.428571

	town_status	urb_v_rur	death_rate_17
0	grown	rural	1.843318
1	grown	rural	0.000000
2	grown	rural	3.830317
3	shrunk	rural	0.000000
4	grown	rural	1.440455

[5 rows x 29 columns]

```
[47]: rate_exp['death_rate_17'].hist(bins=20)
plt.xlabel('Rate (per 10k residents)')
plt.ylabel('Count')
plt.title('Opioid overdose death rate per 10k town residents (2017)')
plt.show()
```




```
[48]: rate_exp.columns
```

```
[48]: Index(['city_death', 'tot_pop_17', 'over_65_count', 'town_actual_2010_pop',  
        'block_est_2010_pop', 'count_error', 'percent_error', 'pop_density',  
        'med_house_inc', 'mean_house_inc', 'less_than_hs_ed', 'over_65_prop',  
        'at_or_below_pov_prop', 'pop_struggling_prop', '2014', '2015', '2016',  
        '2017', '2018', 'pop_change_10_to_17', 'pop_change_10_to_17_est',  
        'pop_change_rate', 'tot_pop_16', 'tot_pop_15', 'tot_pop_14',  
        'tot_pop_13', 'town_status', 'urb_v_rur', 'death_rate_17'],  
        dtype='object')
```

```
[49]: rate_exp_columns = list(rate_exp.columns)  
rate_exp_col_sub = rate_exp_columns[0:1] + rate_exp_columns[7:14] +  
    ↪ rate_exp_columns[21:22] + rate_exp_columns[26:]  
rate_exp_sub = rate_exp[rate_exp_col_sub].copy()  
rate_exp_sub.head()
```

```
[49]:  city_death  pop_density  med_house_inc  mean_house_inc  less_than_hs_ed  \  
0  abington  1932.969130  87156.000000  98809.035505  5.405643  
1  acton     1257.583593  139890.466667  156680.203867  2.456531  
2  acushnet  1152.357871  69624.714286  80333.175842  18.297315  
3  adams     1982.318840  48445.400000  60968.594660  11.862182  
4  agawam    1897.273569  65490.125000  79464.234446  7.748863  
  
    over_65_prop  at_or_below_pov_prop  pop_struggling_prop  pop_change_rate  \  
0      0.151705      0.035754      0.100408      41.428571  
1      0.170582      0.038315      0.041747      218.714286  
2      0.232788      0.040828      0.178406      20.000000  
3      0.214834      0.110854      0.144597     -39.142857  
4      0.223090      0.094819      0.142656      21.142857  
  
    town_status  urb_v_rur  death_rate_17  
0      grown      rural      1.843318  
1      grown      rural      0.000000  
2      grown      rural      3.830317  
3     shrunk      rural      0.000000  
4      grown      rural      1.440455
```

```
[50]: rate_exp_sub_corr = rate_exp_sub.drop(['city_death', 'town_status',  
    ↪ 'urb_v_rur'], axis=1).dropna().corr(method='spearman')  
mask = np.zeros_like(rate_exp_sub_corr, dtype=np.bool)  
mask[np.triu_indices_from(mask)] = True  
f, ax = plt.subplots(figsize=(11, 9))  
cmap = sns.diverging_palette(220, 10, as_cmap=True)  
sns.heatmap(rate_exp_sub_corr, mask=mask, cmap=cmap, vmax=.3, center=0,  
            square=True, linewidths=.5, cbar_kws={"shrink": .5})  
plt.show()
```



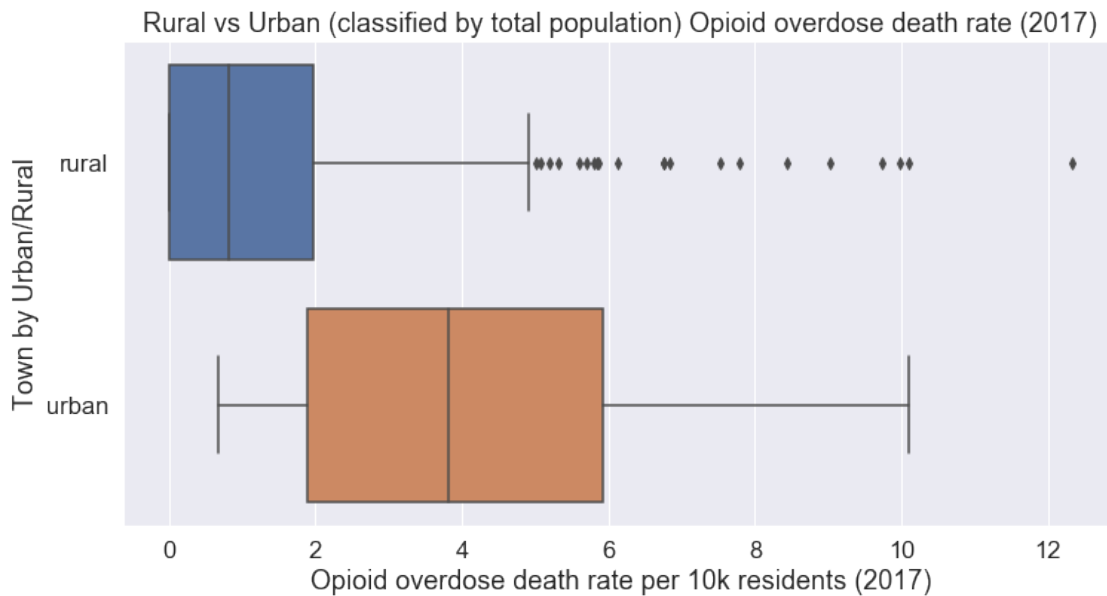
```
[51]: rate_exp_sub_corr.sort_values('death_rate_17', ascending=False)['death_rate_17']
```

```
[51]: death_rate_17      1.000000
      pop_density      0.577145
      pop_change_rate   0.375721
      less_than_hs_ed   0.353322
      at_or_below_pov_prop 0.290039
      pop_struggling_prop 0.270851
      over_65_prop     -0.145125
      med_house_inc     -0.198810
      mean_house_inc    -0.204348
      Name: death_rate_17, dtype: float64
```

Looks like there are some weak positive and negative correlations * pop_density correlation surprisingly high - think it has to do with urban/rural (city vs urban)

```
[52]: plt.figure(figsize=(12, 6))
      sns.boxplot(y='urb_v_rur', x='death_rate_17', orient='h', data=rate_exp_sub)
      plt.ylabel('Town by Urban/Rural')
      plt.xlabel('Opioid overdose death rate per 10k residents (2017)')
```

```
plt.title('Rural vs Urban (classified by total population) Opioid overdose_
→death rate (2017)')
plt.show()
```



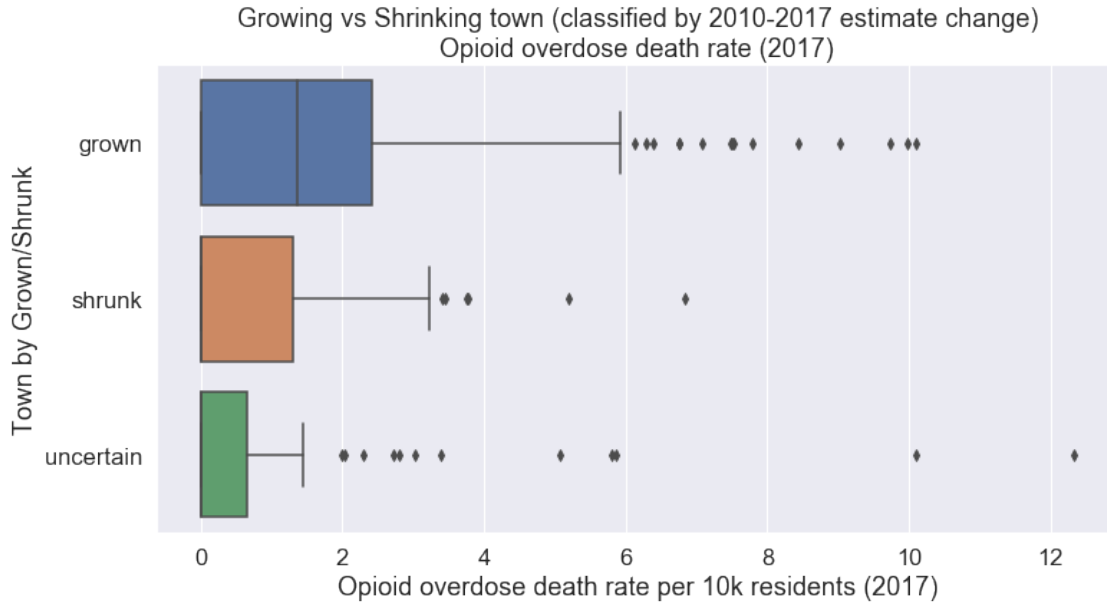
```
[53]: # mean overdose death rate for 2017
print(round(rate_exp_sub['death_rate_17'].mean(), 2))
# standard error of the mean (sem) the rate:
print(round(rate_exp_sub[['death_rate_17']].sem(axis=0), 2))
rate_exp_sub[['urb_v_rur', 'death_rate_17']].groupby('urb_v_rur').mean()
```

```
1.63
death_rate_17    0.12
dtype: float64
```

```
[53]:          death_rate_17
urb_v_rur
rural          1.441318
urban          4.020464
```

“Urban” (high population area) rate seems to be much higher than low population areas

```
[54]: plt.figure(figsize=(12, 6))
sns.boxplot(y='town_status', x='death_rate_17', orient='h', data=rate_exp_sub)
plt.ylabel('Town by Grown/Shrunk')
plt.xlabel('Opioid overdose death rate per 10k residents (2017)')
plt.title('Growing vs Shrinking town (classified by 2010-2017 estimate_
→change)\nOpioid overdose death rate (2017)')
plt.show()
```





```
[55]: rate_exp_sub[['town_status', 'death_rate_17']].groupby('town_status').mean()
```

```
[55]:      death_rate_17
town_status
grown      1.858204
shrunk     1.029335
uncertain  1.006140
```

Small difference with large errors/outliers? - may not be meaningful


```
[56]: plt.figure(figsize=(14, 10))
sns.regplot(x='pop_change_rate', y='death_rate_17',
            data=rate_exp_sub[rate_exp_sub['pop_change_rate'] < 1000])
plt.xlabel('Estimated yearly population change rate')
plt.ylabel('Opioid overdose death rate per 10k residents (2017)')
plt.show()
```



Think the estimated population change will probably not be useful, also most likely these two variables are confounded because population estimates for each year were calculated using the population change rate.


Population density probably related to the urban/rural difference seen earlier - what's the relationship of this variable to the death rate?

```
[57]: plt.figure(figsize=(14, 10))
sns.regplot(x='pop_density', y='death_rate_17', data=rate_exp_sub)
plt.xlabel('Population density (ACS 2017)')
plt.ylabel('Opioid overdose death rate per 10k residents (2017)')
plt.show()
```



Data too noisy to tell if there is a true relationship - this is probably not a useful metric, but what if I try to take the square root of the density (maybe it's too skewed?)

```
[58]: rate_exp_sub['pop_den_sqrt'] = np.sqrt(rate_exp_sub['pop_density'])
plt.figure(figsize=(14, 10))
sns.regplot(x='pop_den_sqrt', y='death_rate_17', data=rate_exp_sub)
plt.xlabel('sqrt(Population density (ACS 2017))')
plt.ylabel('Opioid overdose death rate per 10k residents (2017)')
plt.show()
```



There might be a positive relationship - but still a lot of noise

1.1 Prep for export

Note: I had originally planned to predict/model the death rates instead of the raw counts - but that didn't go well and I switched focus to trying to model the counts instead (but included the total population estimates in the model as a way to account for differences in counts based on population)

```
[59]: full_merge.columns
```

```
[59]: Index(['city_death', 'tot_pop_17', 'over_65_count', 'town_actual_2010_pop',
        'block_est_2010_pop', 'count_error', 'percent_error', 'pop_density',
        'med_house_inc', 'mean_house_inc', 'less_than_hs_ed', 'over_65_prop',
        'at_or_below_pov_prop', 'pop_struggling_prop', '2014', '2015', '2016',
        '2017', '2018', 'pop_change_10_to_17', 'pop_change_10_to_17_est',
        'pop_change_rate', 'tot_pop_16', 'tot_pop_15', 'tot_pop_14',
        'tot_pop_13', 'town_status', 'urb_v_rur'],
        dtype='object')
```

```
[94]: full_merge_cols = list(full_merge.columns)
      full_merge_select = full_merge_cols[0:3] + full_merge_cols[8:19] +
      ↪full_merge_cols[22:]
      print(full_merge_select)
```

```

print(len(full_merge_select))
print(len(set(full_merge_select)))
full_merge_reorder = full_merge_select[:1] + full_merge_select[9:14] +
    →full_merge_select[17:13:-1] + full_merge_select[1:2] + full_merge_select[2:
    →3] + full_merge_select[6:7] + full_merge_select[3:6] + full_merge_select[7:
    →9] + full_merge_select[-1:-3:-1]
print(full_merge_reorder)
print(len(full_merge_reorder))
# make sure no columns were doubl copied
print(len(set(full_merge_reorder)))

```

```

['city_death', 'tot_pop_17', 'over_65_count', 'med_house_inc', 'mean_house_inc',
'less_than_hs_ed', 'over_65_prop', 'at_or_below_pov_prop',
'pop_struggling_prop', '2014', '2015', '2016', '2017', '2018', 'tot_pop_16',
'tot_pop_15', 'tot_pop_14', 'tot_pop_13', 'town_status', 'urb_v_rur']

```

20

20

```

['city_death', '2014', '2015', '2016', '2017', '2018', 'tot_pop_13',
'tot_pop_14', 'tot_pop_15', 'tot_pop_16', 'tot_pop_17', 'over_65_count',
'over_65_prop', 'med_house_inc', 'mean_house_inc', 'less_than_hs_ed',
'at_or_below_pov_prop', 'pop_struggling_prop', 'urb_v_rur', 'town_status']

```

20

20

```

[95]: full_merge_for_csv = full_merge[full_merge_reorder].copy()
full_merge_for_csv.head()

```

```

[95]:  city_death  2014  2015  2016  2017  2018  tot_pop_13  tot_pop_14  \
0  abington      0      6      1      3      5  16109.285714  16150.714286
1    acton       1      2      3      0      1  22580.142857  22798.857143
2  acushnet      0      4      2      4      0  10363.000000  10383.000000
3    adams       2      3      1      0      4   8367.571429   8328.428571
4  agawam        1      2      0      4      8  27684.428571  27705.571429

      tot_pop_15  tot_pop_16  tot_pop_17  over_65_count  over_65_prop  \
0  16192.142857  16233.571429      16275          2469      0.151705
1  23017.571429  23236.285714      23455          4001      0.170582
2  10403.000000  10423.000000      10443          2431      0.232788
3   8289.285714   8250.142857       8211          1764      0.214834
4  27726.714286  27747.857143      27769          6195      0.223090

      med_house_inc  mean_house_inc  less_than_hs_ed  at_or_below_pov_prop  \
0   87156.000000    98809.035505          5.405643          0.035754
1  139890.466667   156680.203867          2.456531          0.038315
2   69624.714286   80333.175842         18.297315          0.040828
3   48445.400000   60968.594660         11.862182          0.110854
4   65490.125000   79464.234446          7.748863          0.094819

```

	pop_struggling_prop	urb_v_rur	town_status
0	0.100408	rural	grown
1	0.041747	rural	grown
2	0.178406	rural	grown
3	0.144597	rural	shrunk
4	0.142656	rural	grown

```
[96]: # final file:
#full_merge_for_csv.to_csv("../data/tidy_data/overdose_death_count_acs_merge.
→csv", index=False)
```

Old files derived from this notebook (to keep track to reorganize other notebooks - can ignore otherwise)

```
[ ]: #full_merge.to_csv("../data/tidy_data/
→death_count_norm_to_pop_and_acs_town_demographics_merge_all_cols.csv",
→index=False)
#full_merge_for_csv.to_csv("../data/tidy_data/
→death_count_norm_to_pop_and_acs_town_demographics_merge.csv", index=False)
```